





Lecture:

Face Recognition and Feature Reduction

Juan Carlos Niebles and Ranjay Krishna
Stanford Vision and Learning Lab

Recap - Curse of dimensionality

- Assume 5000 points uniformly distributed in the unit hypercube and we want to apply 5-NN. Suppose our query point is at the origin.
 - In 1-dimension, we must go a distance of $5/5000=0.001$ on the average to capture 5 nearest neighbors.
 - In 2 dimensions, we must go $\sqrt{0.001}$ to get a square that contains 0.001 of the volume.
 - In d dimensions, we must go $(0.001)^{1/d}$



What we will learn today

- Singular value decomposition
- Principal Component Analysis (PCA)
- Image compression

What we will learn today

- Singular value decomposition
- Principal Component Analysis (PCA)
- Image compression

Singular Value Decomposition (SVD)

- There are several computer algorithms that can “factorize” a matrix, representing it as the product of some other matrices
- The most useful of these is the Singular Value Decomposition.
- Represents any matrix \mathbf{A} as a product of three matrices: $\mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$
- Python command:
 - $[\mathbf{U}, \mathbf{S}, \mathbf{V}] = \text{numpy.linalg.svd}(\mathbf{A})$

Singular Value Decomposition (SVD)

$$\mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \mathbf{A}$$

- Where \mathbf{U} and \mathbf{V} are rotation matrices, and $\mathbf{\Sigma}$ is a scaling matrix. For example:

$$\begin{matrix} U & & \Sigma & & V^T & & A \\ \begin{bmatrix} -.40 & .916 \\ .916 & .40 \end{bmatrix} & \times & \begin{bmatrix} 5.39 & 0 \\ 0 & 3.154 \end{bmatrix} & \times & \begin{bmatrix} -.05 & .999 \\ .999 & .05 \end{bmatrix} & = & \begin{bmatrix} 3 & -2 \\ 1 & 5 \end{bmatrix} \end{matrix}$$

Singular Value Decomposition (SVD)

- Beyond 2x2 matrices:
 - In general, if \mathbf{A} is $m \times n$, then \mathbf{U} will be $m \times m$, $\mathbf{\Sigma}$ will be $m \times n$, and \mathbf{V}^T will be $n \times n$.
 - (Note the dimensions work out to produce $m \times n$ after multiplication)

$$\begin{matrix} U & & \Sigma & & V^T & & A \\ \begin{bmatrix} -.39 & -.92 \\ -.92 & .39 \end{bmatrix} & \times & \begin{bmatrix} 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{bmatrix} & \times & \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} & = & \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \end{matrix}$$

Singular Value Decomposition (SVD)

- **U** and **V** are always rotation matrices.
 - Geometric rotation may not be an applicable concept, depending on the matrix. So we call them “unitary” matrices – each column is a unit vector.
- **Σ** is a diagonal matrix
 - The number of nonzero entries = rank of **A**
 - The algorithm always sorts the entries high to low

$$\begin{matrix} U \\ \begin{bmatrix} -.39 & -.92 \\ -.92 & .39 \end{bmatrix} \end{matrix} \times \begin{matrix} \Sigma \\ \begin{bmatrix} 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{bmatrix} \end{matrix} \times \begin{matrix} V^T \\ \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{matrix} = \begin{matrix} A \\ \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \end{matrix}$$

SVD Applications

- We've discussed SVD in terms of geometric transformation matrices
- But SVD of an image matrix can also be very useful
- To understand this, we'll look at a less geometric interpretation of what SVD is doing

SVD Applications

$$\overset{U}{\begin{bmatrix} -.39 & -.92 \\ -.92 & .39 \end{bmatrix}} \times \overset{\Sigma}{\begin{bmatrix} 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{bmatrix}} \times \overset{V^T}{\begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix}} = \overset{A}{\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}}$$

- Look at how the multiplication works out, left to right:
- Column 1 of **U** gets scaled by the first value from **Σ**.

$$\overset{U\Sigma}{\begin{bmatrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{bmatrix}} \times \overset{V^T}{\begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix}} = \overset{A_{\text{partial}}}{\begin{bmatrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{bmatrix}}$$

- The resulting vector gets scaled by row 1 of **V^T** to produce a contribution to the columns of **A**

SVD Applications

decomposition USigma

$$\begin{aligned}
 & \begin{matrix} U\Sigma \\ \begin{bmatrix} -3.67 & -0.71 & 0 \\ -8.8 & 0.30 & 0 \end{bmatrix} \end{matrix} \times \begin{matrix} V^T \\ \begin{bmatrix} -0.42 & -0.57 & -0.70 \\ 0.81 & 0.11 & -0.58 \\ 0.41 & -0.82 & 0.41 \end{bmatrix} \end{matrix} = \begin{matrix} A_{\text{partial}} \\ \begin{bmatrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{bmatrix} \end{matrix} \\
 + & \begin{matrix} U\Sigma \\ \begin{bmatrix} -3.67 & -0.71 & 0 \\ -8.8 & 0.30 & 0 \end{bmatrix} \end{matrix} \times \begin{matrix} V^T \\ \begin{bmatrix} -0.42 & -0.57 & -0.70 \\ 0.81 & 0.11 & -0.58 \\ 0.41 & -0.82 & 0.41 \end{bmatrix} \end{matrix} = \begin{matrix} A_{\text{partial}} \\ \begin{bmatrix} -0.6 & -0.1 & 0.4 \\ 0.2 & 0 & -0.2 \end{bmatrix} \end{matrix} \\
 = & \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}
 \end{aligned}$$

- Each product of (column i of U) \cdot (value i from Σ) \cdot (row i of V^T) produces a component of the final A .

SVD Applications

$$\begin{array}{c}
 \begin{array}{c} U\Sigma \\ \begin{bmatrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{bmatrix} \end{array} \times \begin{array}{c} V^T \\ \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{array} \quad \begin{array}{c} A_{\text{partial}} \\ \begin{bmatrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{bmatrix} \end{array} \quad \begin{array}{c} A \\ \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \end{array} \\
 \\
 \begin{array}{c} U\Sigma \\ \begin{bmatrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{bmatrix} \end{array} \times \begin{array}{c} V^T \\ \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{array} \quad \begin{array}{c} A_{\text{partial}} \\ \begin{bmatrix} -.6 & -.1 & .4 \\ .2 & 0 & -.2 \end{bmatrix} \end{array}
 \end{array}$$

- We're building **A** as a linear combination of the columns of **U**
- Using all columns of **U**, we'll rebuild the original matrix perfectly
- But, in real-world data, often we can just use the first few columns of **U** and we'll get something close (e.g. the first **A_{partial}**, above)

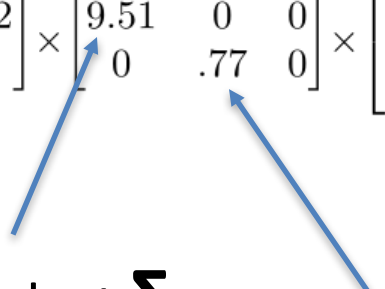
SVD Applications

$$\begin{array}{c}
 U\Sigma \\
 \begin{bmatrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{bmatrix} \times \begin{array}{c} V^T \\ \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{array} \begin{array}{c} A_{\text{partial}} \\ \begin{bmatrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{bmatrix} \end{array}
 \end{array}
 \begin{array}{c}
 A \\
 \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}
 \end{array}$$

$$\begin{array}{c}
 U\Sigma \\
 \begin{bmatrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{bmatrix} \times \begin{array}{c} V^T \\ \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{array} \begin{array}{c} A_{\text{partial}} \\ \begin{bmatrix} -.6 & -.1 & .4 \\ .2 & 0 & -.2 \end{bmatrix} \end{array}
 \end{array}$$

- We can call those first few columns of \mathbf{U} the *Principal Components* of the data
- They show the major patterns that can be added to produce the columns of the original matrix
- The rows of \mathbf{V}^T show how the *principal components* are mixed to produce the columns of the matrix


SVD Applications

$$\begin{matrix} U \\ \begin{bmatrix} -.39 & -.92 \\ -.92 & .39 \end{bmatrix} \end{matrix} \times \begin{matrix} \Sigma \\ \begin{bmatrix} 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{bmatrix} \end{matrix} \times \begin{matrix} V^T \\ \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{matrix} = \begin{matrix} A \\ \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \end{matrix}$$


We can look at Σ to see that the first column has a large effect

while the second column has a much smaller effect in this example

SVD Applications



- For this image, using **only the first 10** of 300 principal components produces a recognizable reconstruction U matrix
- So, SVD can be used for image compression

SVD for symmetric matrices

- If A is a symmetric matrix, it can be decomposed as the following:

$$A = \Phi \Sigma \Phi^T$$

- Compared to a traditional SVD decomposition, ~~$U = V^T$~~ and is an orthogonal matrix.
 $U = V$

Principal Component Analysis

$$\begin{array}{c} U\Sigma \\ \begin{bmatrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{bmatrix} \end{array} \times \begin{array}{c} V^T \\ \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{array} = \begin{array}{c} A_{\text{partial}} \\ \begin{bmatrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{bmatrix} \end{array}$$

- Remember, columns of \mathbf{U} are the Principal Components of the data: the major patterns that can be added to produce the columns of the original matrix
- One use of this is to construct a matrix where each column is a separate data sample
- Run SVD on that matrix, and look at the first few columns of \mathbf{U} to see patterns that are common among the columns
- This is called *Principal Component Analysis* (or PCA) of the data samples

Principal Component Analysis

$$\begin{matrix} U\Sigma \\ \begin{bmatrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{bmatrix} \end{matrix} \times \begin{matrix} V^T \\ \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{matrix} = \begin{matrix} A_{\text{partial}} \\ \begin{bmatrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{bmatrix} \end{matrix}$$

- Often, raw data samples have a lot of redundancy and patterns
- PCA can allow you to represent data samples as weights on the principal components, rather than using the original raw form of the data
- By representing each sample as just those weights, you can represent just the “meat” of what’s different between samples.
- This minimal representation makes machine learning and other algorithms much more efficient

How is SVD computed?

- For this class: tell PYTHON to do it. Use the result.
- But, if you're interested, one computer algorithm to do it makes use of Eigenvectors!

Eigenvector definition

- Suppose we have a square matrix **A**. We can solve for vector x and scalar λ such that $Ax = \lambda x$
- In other words, find vectors where, if we transform them with **A**, the only effect is to scale them with no change in direction.
- These vectors are called eigenvectors (German for “self vector” of the matrix), and the scaling factors λ are called eigenvalues
- An $m \times m$ matrix will have $\leq m$ eigenvectors where λ is nonzero

Finding eigenvectors

- Computers can find an x such that $Ax = \lambda x$ using this iterative algorithm:
 - X = random unit vector
 - while(x hasn't converged)
 - $X = Ax$
 - normalize x
- x will quickly converge to an eigenvector
- Some simple modifications will let this algorithm find all eigenvectors

Finding SVD

- Eigenvectors are for square matrices, but SVD is for all matrices
Gram matrix : $V^T * V$
- To do $\text{svd}(A)$, computers can do this:
 - Take eigenvectors of AA^T (matrix is always square).
 - These eigenvectors are the **columns** of **U** .
 - Square root of eigenvalues are the singular values (the entries of Σ).
 - Take eigenvectors of $A^T A$ (matrix is always square).
 - These eigenvectors are **columns** of **V** (or rows of V^T)

assumed A (- $Rm * n$)
 Then U (- $Rm * m$)
 V (- $Rn * n$)
 This way, AA^T (- $Rm * m$)
 $A^T A$ (- $Rn * n$)

Finding SVD

- Moral of the story: SVD is fast, even for large matrices
- It's useful for a lot of stuff
- There are also other algorithms to compute SVD or part of the SVD
 - Python's `np.linalg.svd()` command has options to efficiently compute only what you need, if performance becomes an issue

A detailed geometric explanation of SVD is here:

<http://www.ams.org/samplings/feature-column/fcarc-svd>

What we will learn today

- Introduction to face recognition
- **Principal Component Analysis (PCA)**
- Image compression

Covariance

- Variance and Covariance are a measure of the “spread” of a set of points around their center of mass (mean)
- Variance – measure of the deviation from the mean for points in one dimension e.g. heights
- Covariance as a measure of how much each of the dimensions vary from the mean with respect to each other.
- Covariance is measured between 2 dimensions to see if there is a relationship between the 2 dimensions e.g. number of hours studied & marks obtained.
- The covariance between one dimension and itself is the variance

Covariance

$$\text{covariance}(X,Y) = \frac{\sum_{i=1}^n (\bar{X}_i - \bar{X})(\bar{Y}_i - \bar{Y})}{(n-1)}$$

$$\text{cov}(X, Y) = E\{[X - E(X)][Y - E(Y)]\}$$

$$\text{cov}(X, Y) = E(XY) - E(X)E(Y)$$

- So, if you had a 3-dimensional data set (x,y,z), then you could measure the covariance between the x and y dimensions, the y and z dimensions, and the x and z dimensions. Measuring the covariance between x and x, or y and y, or z and z would give you the variance of the x, y and z dimensions respectively

Covariance matrix

- Representing Covariance between dimensions as a matrix e.g. for 3 dimensions

$$C = \begin{bmatrix} \text{cov}(x,x) & \text{cov}(x,y) & \text{cov}(x,z) \\ \text{cov}(y,x) & \text{cov}(y,y) & \text{cov}(y,z) \\ \text{cov}(z,x) & \text{cov}(z,y) & \text{cov}(z,z) \end{bmatrix}$$


Variances

- Diagonal is the variances of x, y and z
- cov(x,y) = cov(y,x) hence matrix is symmetrical about the diagonal
- N-dimensional data will result in NxN covariance matrix

Covariance

- What is the interpretation of covariance calculations?
 - e.g.: 2 dimensional data set
 - x : number of hours studied for a subject
 - y : marks obtained in that subject
 - covariance value is say: 104.53
 - what does this value mean?

Covariance interpretation




Covariance interpretation


- Exact value is not as important as it's sign.
- A **positive value** of covariance indicates both dimensions increase or decrease together e.g. as the number of hours studied increases, the marks in that subject increase.
- A **negative value** indicates while one increases the other decreases, or vice-versa e.g. active social life at PSU vs performance in CS dept.
- If **covariance is zero**: the two dimensions are independent of each other e.g. heights of students vs the marks obtained in a subject

Example data

Covariance between the two axis is high. Can we reduce the number of dimensions to just 1?




Geometric interpretation of PCA



Geometric interpretation of PCA

- Let's say we have a set of 2D data points x . But we see that all the points lie on a line in 2D.
- So, 2 dimensions are redundant to express the data. We can express all the points with just one dimension.





PCA: Principle Component Analysis

- Given a set of points, how do we know if they can be compressed like in the previous example?
 - The answer is to look into the correlation between the points
 - The tool for doing this is called PCA

PCA Formulation

- Basic idea:
 - If the data lives in a subspace, it is going to look very flat when viewed from the full space, e.g.



Slide inspired by N. Vasconcelos

PCA Formulation


- Assume \mathbf{x} is Gaussian with covariance Σ .
- Recall that a gaussian is defined with it's mean and variance:

$$\mathbf{X} \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma)$$

- Recall that $\boldsymbol{\mu}$ and Σ of a gaussian are defined as:

$$\boldsymbol{\mu} = \mathbf{E}[\mathbf{X}] = [\mathbf{E}[X_1], \mathbf{E}[X_2], \dots, \mathbf{E}[X_k]]^T$$

$$\Sigma =: \mathbf{E}[(\mathbf{X} - \boldsymbol{\mu})(\mathbf{X} - \boldsymbol{\mu})^T] = [\text{Cov}[X_i, X_j]; 1 \leq i, j \leq k]$$



PCA formulation

- Since gaussians are symmetric, it's covariance matrix is also a symmetric matrix. So we can express it as:

$$-\Sigma = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T = \mathbf{U}\mathbf{\Lambda}^{1/2}(\mathbf{U}\mathbf{\Lambda}^{1/2})^T$$

$$\mathbf{X} \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma) \iff \mathbf{X} \sim \boldsymbol{\mu} + \mathbf{U}\mathbf{\Lambda}^{1/2}\mathcal{N}(0, \mathbf{I})$$

$$\iff \mathbf{X} \sim \boldsymbol{\mu} + \mathbf{U}\mathcal{N}(0, \mathbf{\Lambda}).$$


$$\mathbf{U}^T * (\mathbf{X} - \boldsymbol{\mu}) \sim \mathcal{N}(0, \mathbf{\Lambda})$$

$$\mathbf{U}^T * \mathbf{X} \sim \mathcal{N}(\mathbf{U}^T * \boldsymbol{\mu}, \mathbf{\Lambda})$$

PCA Formulation

- If x is Gaussian with covariance Σ ,

- Principal components ϕ_i are the eigenvectors of Σ
- Principal lengths λ_i are the eigenvalues of Σ



- by computing the eigenvalues we know the data is
 - Not flat if $\lambda_1 \approx \lambda_2$
 - Flat if $\lambda_1 \gg \lambda_2$

Slide inspired by N. Vasconcelos

PCA Algorithm (training)

► Given sample $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$, $x_i \in \mathcal{R}^d$

- compute sample mean: $\hat{\mu} = \frac{1}{n} \sum_i (\mathbf{x}_i)$

- compute sample covariance: $\hat{\Sigma} = \frac{1}{n} \sum_i (\mathbf{x}_i - \hat{\mu})(\mathbf{x}_i - \hat{\mu})^T$

- compute eigenvalues and eigenvectors of $\hat{\Sigma}$

$$\hat{\Sigma} = \Phi \Lambda \Phi^T, \quad \Lambda = \text{diag}(\sigma_1^2, \dots, \sigma_n^2) \quad \Phi^T \Phi = I$$

- order eigenvalues $\sigma_1^2 > \dots > \sigma_n^2$

- if, for a certain k , $\sigma_k \ll \sigma_1$ eliminate the eigenvalues and eigenvectors above k .

Slide inspired by N. Vasconcelos

PCA Algorithm (testing)

- Given principal components $\phi_i, i \in 1, \dots, k$ and a test sample $\mathcal{T} = \{\mathbf{t}_1, \dots, \mathbf{t}_n\}, \mathbf{t}_i \in \mathcal{R}^d$

- subtract mean to each point $\mathbf{t}'_i = \mathbf{t}_i - \hat{\mu}$

- project onto eigenvector space $\mathbf{y}_i = \mathbf{A}\mathbf{t}'_i$ where

$$\mathbf{A} = \begin{bmatrix} \phi_1^T \\ \vdots \\ \phi_k^T \end{bmatrix}$$

- use $\mathcal{T}' = \{\mathbf{y}_1, \dots, \mathbf{y}_n\}$ to estimate class conditional densities and do all further processing on \mathbf{y} .

Slide inspired by N. Vasconcelos

PCA by SVD

- An alternative manner to compute the principal components, based on singular value decomposition
- Quick reminder: SVD
 - Any real $n \times m$ matrix ($n > m$) can be decomposed as

$$A = M \Pi N^T$$

reduced singular value
decomposition

- Where M is an $(n \times m)$ column orthonormal matrix of left singular vectors (columns of M)
- Π is an $(m \times m)$ diagonal matrix of singular values
- N^T is an $(m \times m)$ row orthonormal matrix of right singular vectors (columns of N)

$$M^T M = I \quad N^T N = I$$

Slide inspired by N. Vasconcelos

PCA by SVD

- To relate this to PCA, we consider the data matrix

$$X = \begin{bmatrix} | & & | \\ x_1 & \dots & x_n \\ | & & | \end{bmatrix}$$

- The sample mean is

$$\mu = \frac{1}{n} \sum_i x_i = \frac{1}{n} \begin{bmatrix} | & & | \\ x_1 & \dots & x_n \\ | & & | \end{bmatrix} \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix} = \frac{1}{n} X \mathbf{1}$$

Slide inspired by N. Vasconcelos

PCA by SVD

- Center the data by subtracting the mean to each column of X
- The centered data matrix is

$$X_c = \begin{bmatrix} | & & | \\ X_1 & \dots & X_n \\ | & & | \end{bmatrix} - \begin{bmatrix} | & & | \\ \mu & \dots & \mu \\ | & & | \end{bmatrix}$$
$$= X - \mu 1^T = X - \frac{1}{n} X 1 1^T = X \left(I - \frac{1}{n} 1 1^T \right)$$

Slide inspired by N. Vasconcelos

PCA by SVD

- The sample covariance matrix is

$$\Sigma = \frac{1}{n} \sum_i (x_i - \mu)(x_i - \mu)^T = \frac{1}{n} \sum_i x_i^c (x_i^c)^T$$

where x_i^c is the i^{th} column of X_c

- This can be written as

$$\Sigma = \frac{1}{n} \begin{bmatrix} | & & | \\ x_1^c & \dots & x_n^c \\ | & & | \end{bmatrix} \begin{bmatrix} - & x_1^c & - \\ \vdots & \vdots & \vdots \\ - & x_n^c & - \end{bmatrix} = \frac{1}{n} X_c X_c^T$$

Slide inspired by N. Vasconcelos

PCA by SVD

- The matrix

$$X_c^T = \begin{bmatrix} - & x_1^c & - \\ & \vdots & \\ - & x_n^c & - \end{bmatrix}$$

in form of reduced svd
 Π is diagonal

is real ($n \times d$). Assuming $n > d$ it has SVD decomposition

$$X_c^T = M \Pi N^T$$

$$M^T M = I \quad N^T N = I$$

and

$$\Sigma = \frac{1}{n} X_c X_c^T = \frac{1}{n} N \Pi M^T M \Pi N^T = \frac{1}{n} N \Pi^2 N^T$$

Slide inspired by N. Vasconcelos

PCA by SVD

$$\Sigma = \mathbf{N} \left(\frac{1}{n} \mathbf{\Pi}^2 \right) \mathbf{N}^T$$

- Note that \mathbf{N} is $(d \times d)$ and orthonormal, and $\mathbf{\Pi}^2$ is diagonal. This is just the eigenvalue decomposition of Σ
- It follows that
 - The eigenvectors of Σ are the columns of \mathbf{N}
 - The eigenvalues of Σ are

$$\lambda_i = \frac{1}{n} \pi_i^2$$

- This gives an alternative algorithm for PCA

Slide inspired by N. Vasconcelos

PCA by SVD

- In summary, computation of PCA by SVD
- Given X with one example per column
 - Create the centered data matrix

$$X_c^T = \left(I - \frac{1}{n} \mathbf{1} \mathbf{1}^T \right) X^T$$

- Compute its SVD

$$X_c^T = M \Pi N^T$$


- Principal components are columns of N , eigenvalues are

$$\lambda_i = \frac{1}{n} \pi_i^2$$

Slide inspired by N. Vasconcelos

Rule of thumb for finding the number of PCA components

- A natural measure is to pick the eigenvectors that explain p% of the data variability
 - Can be done by plotting the ratio r_k as a function of k



$$r_k = \frac{\sum_{i=1}^k \lambda_i^2}{\sum_{i=1}^n \lambda_i^2}$$

- E.g. we need 3 eigenvectors to cover 70% of the variability of this dataset

Slide inspired by N. Vasconcelos

What we will learn today


- Introduction to face recognition
- Principal Component Analysis (PCA)
- Image compression

Original Image



- Divide the original 372x492 image into patches:
 - Each patch is an instance that contains 12x12 pixels on a grid
- View each as a 144-D vector

L_2 error and PCA dim




PCA compression: 144D \rightarrow 60D



PCA compression: 144D \rightarrow 16D




16 most important eigenvectors



PCA compression: 144D \rightarrow 6D






6 most important eigenvectors




PCA compression: 144D \rightarrow 3D



3 most important eigenvectors



PCA compression: 144D \rightarrow 1D



What we have learned today

- Introduction to face recognition
- Principal Component Analysis (PCA)
- Image compression